

# Task1

February 7, 2018

## 1 Task1

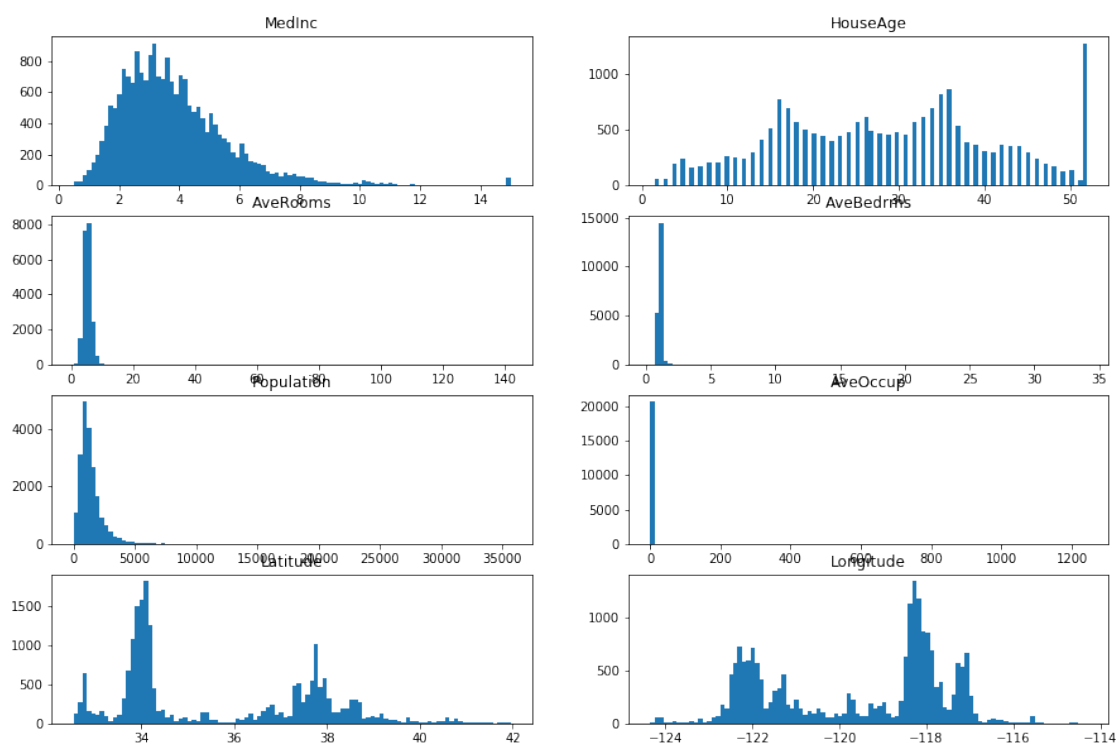

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import sklearn.datasets
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

### 1.1

```
In [18]: california_dataset = sklearn.datasets.fetch_california_housing()

plt.hist(california_dataset['target'], bins=100)
plt.grid(color='gray', linestyle='-', linewidth=0.5, alpha=0.3)
plt.title('target')
plt.show()

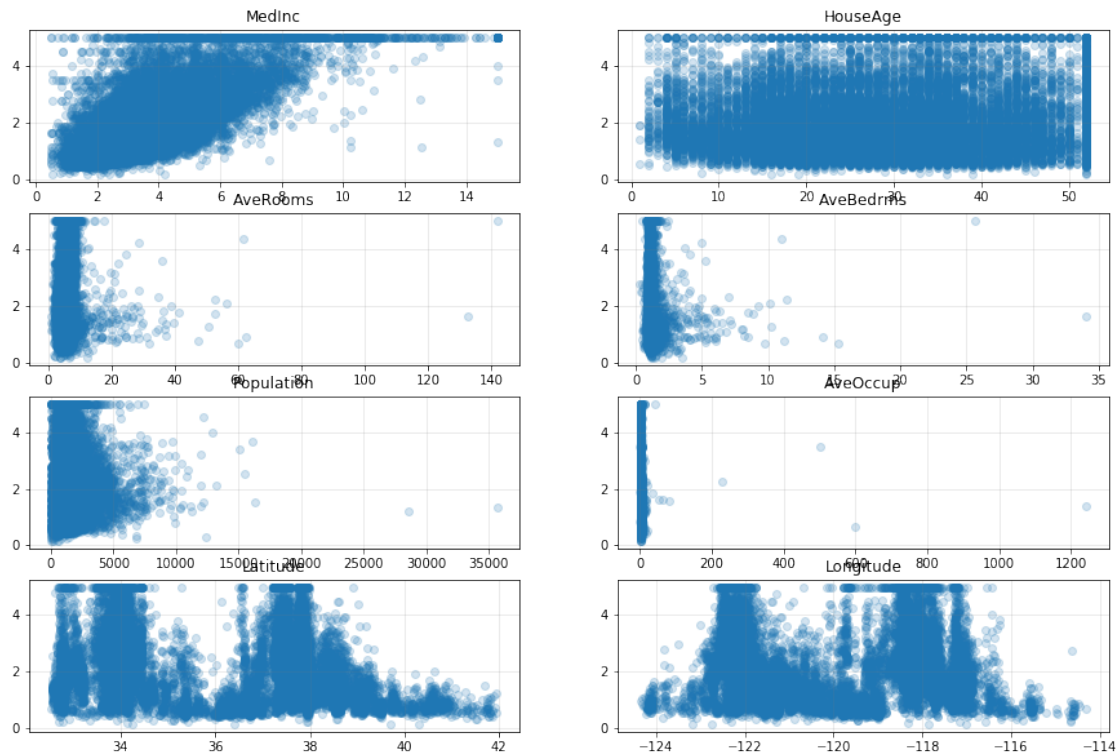
fig, axes = plt.subplots(4, 2, figsize=(15, 10))
for i in range(0, 4):
    for j in range(0, 2):
        axes[i, j].hist(california_dataset['data'][:, i*2+j], bins=100)
        axes[i, j].set_title(california_dataset['feature_names'][i*2+j])
plt.show()
```



From the above eight feature plots, 'AveRooms', 'AveBedrms', 'Population' and 'AveOccup' have inappropriate x scale, which means there are some outliers on these four features. We may need to remove these outliers, than re-scale the x axis.

## 1.2

```
In [4]: fig, axes = plt.subplots(4,2,figsize=(15,10))
        for i in range(0,4):
            for j in range(0,2):
                axes[i,j].scatter(california_dataset['data'][:,i*2+j],california_dataset['target'])
                axes[i,j].set_title(california_dataset['feature_names'][i*2+j])
        plt.show()
```



## 1.3

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(california_dataset['data'],california_dataset['target'])
        LRScore=np.mean(cross_val_score(LinearRegression(), X_train, y_train))
        RidgeScore=np.mean(cross_val_score(Ridge(), X_train, y_train))
        LassoScore=np.mean(cross_val_score(Lasso(), X_train, y_train))
        ElasticNetScore=np.mean(cross_val_score(ElasticNet(), X_train, y_train))
        print('LRScore:{}\nRidgeScore:{}\nLassoScore:{}\nElasticNetScore:{}'.format(LRScore,RidgeScore,LassoScore,ElasticNetScore))
```

LRScore:0.6102422922538064

RidgeScore:0.6102437749800811

```
LassoScore:0.2819718569982819
ElasticNetScore:0.4229514321435757
```

Scaling the features

```
In [13]: scaler = StandardScaler()
        scaler.fit(X_train)
        X_train_scaled = scaler.transform(X_train)
        LRScore_S=np.mean(cross_val_score(LinearRegression(), X_train_scaled, y_train))
        RidgeScore_S=np.mean(cross_val_score(Ridge(), X_train_scaled, y_train))
        LassoScore_S=np.mean(cross_val_score(Lasso(), X_train_scaled, y_train))
        ElasticNetScore_S=np.mean(cross_val_score(ElasticNet(), X_train_scaled, y_train))
        print('LRScore:{}\nRidgeScore:{}\nLassoScore:{}\nElasticNetScore:{}'.format(LRScore_S,
        RidgeScore_S, LassoScore_S, ElasticNetScore_S))

LRScore:0.6102422922538069
RidgeScore:0.6102441875598729
LassoScore:-0.00015307833209199373
ElasticNetScore:0.20448945641200397
```

According to the results, Scaling doesn't help the OLR and Ridge. Besides, it even has a negative influence on the Lasso and ElasticNet.

## 1.4

```
In [20]: # Ridge
        param_grid = {'alpha': np.logspace(-3, 3, 13)}
        grid = GridSearchCV(Ridge(), param_grid, return_train_score=True)
        grid.fit(X_train, y_train)
        plt.plot(param_grid['alpha'], grid.cv_results_['mean_train_score'], c='blue', label='mean_train_score')
        plt.plot(param_grid['alpha'], grid.cv_results_['mean_test_score'], c='red', label='mean_test_score')
        plt.xlabel('alpha')
        plt.ylabel('mean cv score')
        plt.xscale('log')
        plt.legend()
        plt.show()
        ridge = grid.best_estimator_
        print('best score:{}\nbest parameters:{}'.format(grid.best_score_, grid.best_params_))

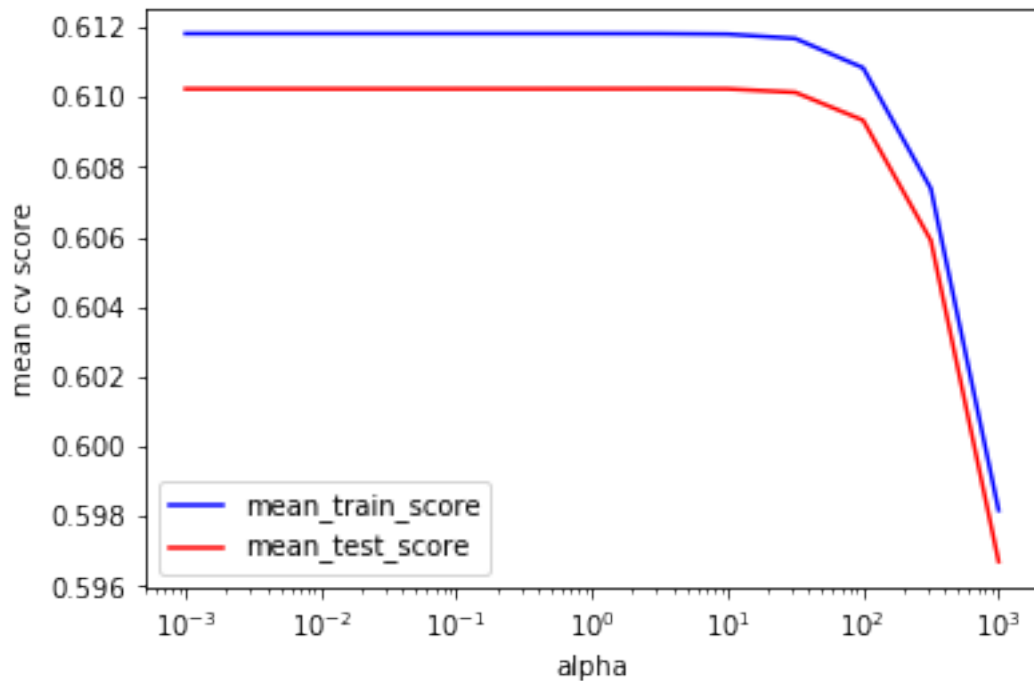
#Lasso
        param_grid = {'alpha': np.logspace(-3, 0, 13)}
        grid = GridSearchCV(Lasso(), param_grid, return_train_score=True)
        grid.fit(X_train, y_train)
        plt.plot(param_grid['alpha'], grid.cv_results_['mean_train_score'], c='blue', label='mean_train_score')
        plt.plot(param_grid['alpha'], grid.cv_results_['mean_test_score'], c='red', label='mean_test_score')
        plt.xlabel('alpha')
        plt.ylabel('mean cv score')
```

```

plt.xscale('log')
plt.legend()
plt.show()
lasso = grid.best_estimator_
print('best score:{}'.format(grid.best_score_))

# ElasticNet
param_grid = {'alpha': np.logspace(-3, 2, 10), 'l1_ratio': [0.01, .1, .5, .9, .98, 1]}
grid = GridSearchCV(ElasticNet(), param_grid, return_train_score=True)
grid.fit(X_train, y_train)
res = pd.pivot_table(pd.DataFrame(grid.cv_results_), values='mean_test_score', index='alpha')
plt.imshow(res, extent=[0, 1, 0, 100], aspect="auto")
plt.colorbar()
plt.show()
en = grid.best_estimator_
print('best score:{}'.format(grid.best_score_))

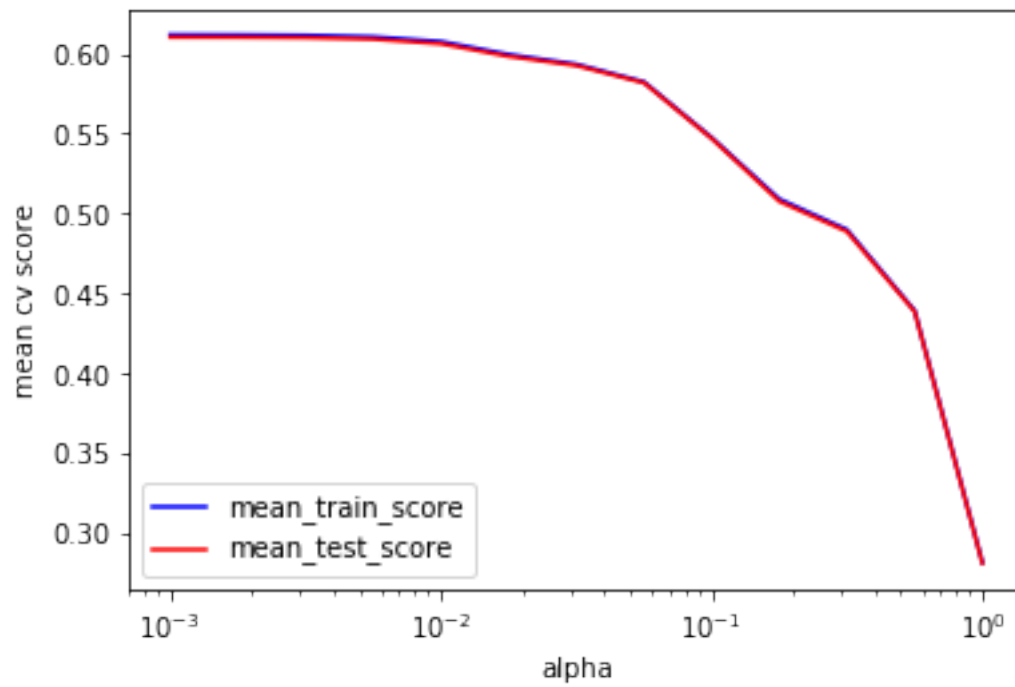
```




```

best score:0.6102457318944084
best parameters: {'alpha': 3.1622776601683795}

```



best score:0.6102299335181556  
best parameters: {'alpha': 0.001}



```
best score:0.6102407928395485
best parameters: {'alpha': 0.001, 'l1_ratio': 0.01}
```

According to the results, GridSearch only can improve Lasso and ElasticNet.

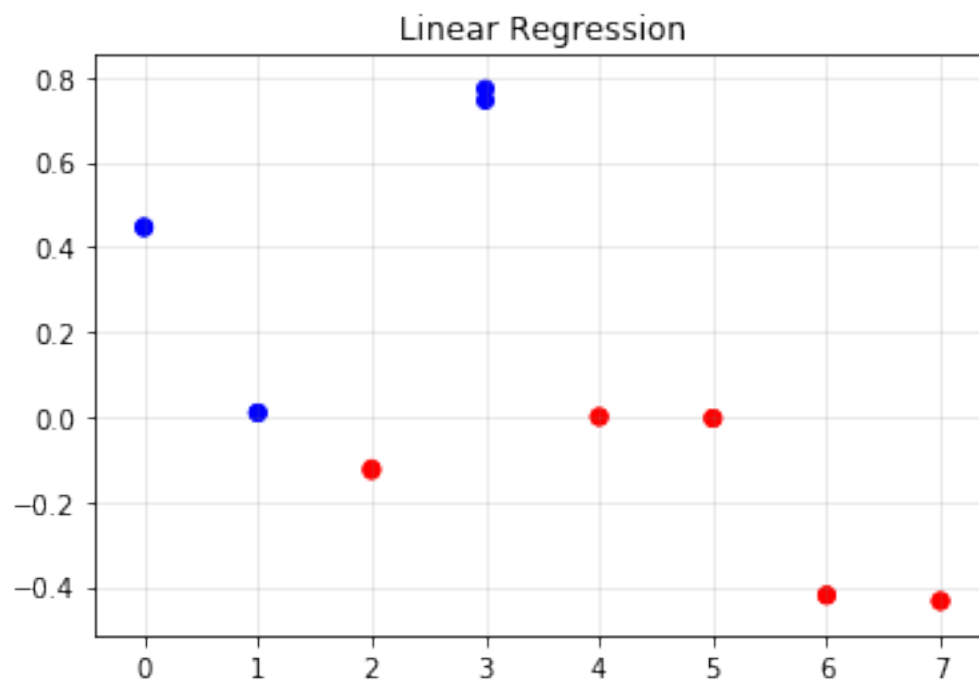
## 1.5

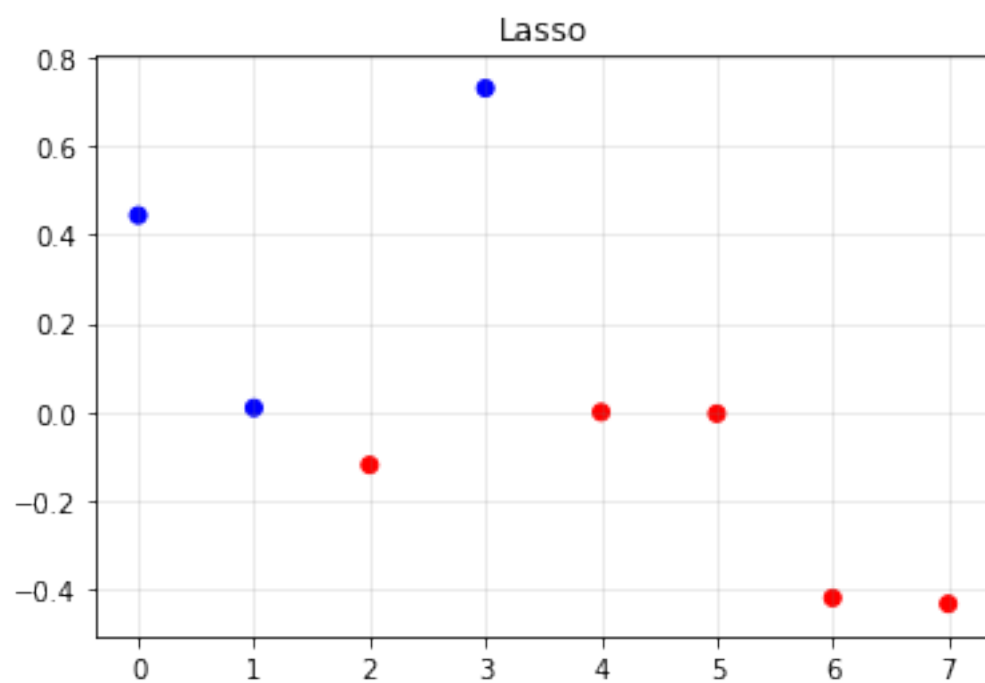
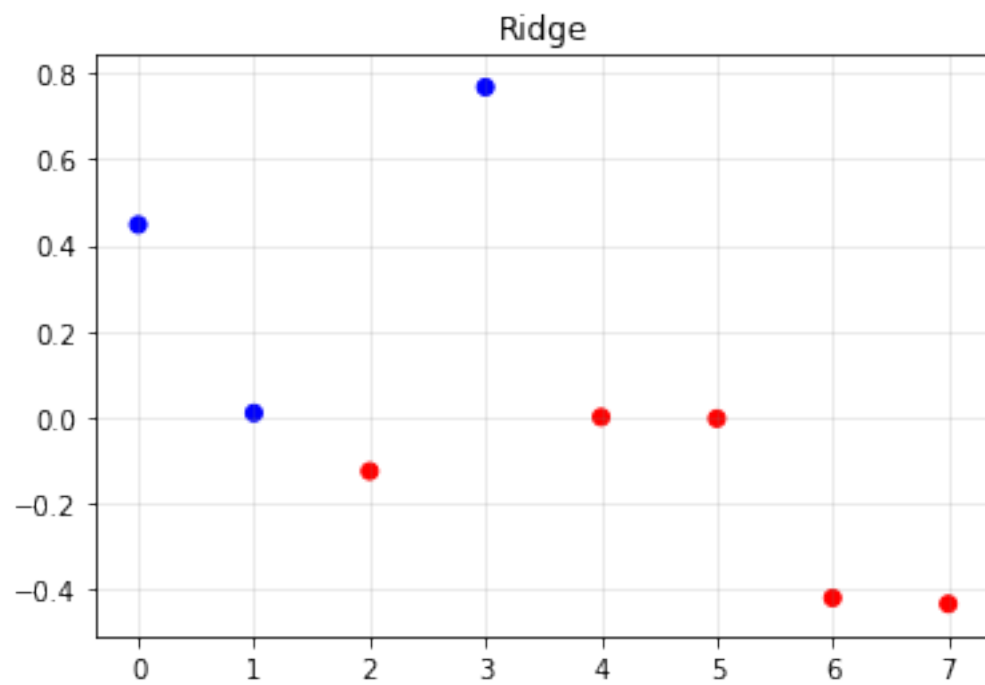
```
In [17]: lr = LinearRegression().fit(X_train, y_train)
plt.scatter(range(X_train.shape[1]), lr.coef_, c=np.sign(lr.coef_), cmap="bwr_r")
plt.title('Linear Regression')
plt.show()

plt.scatter(range(X_train.shape[1]), ridge.coef_, c=np.sign(ridge.coef_), cmap="bwr_r")
plt.title('Ridge')
plt.show()

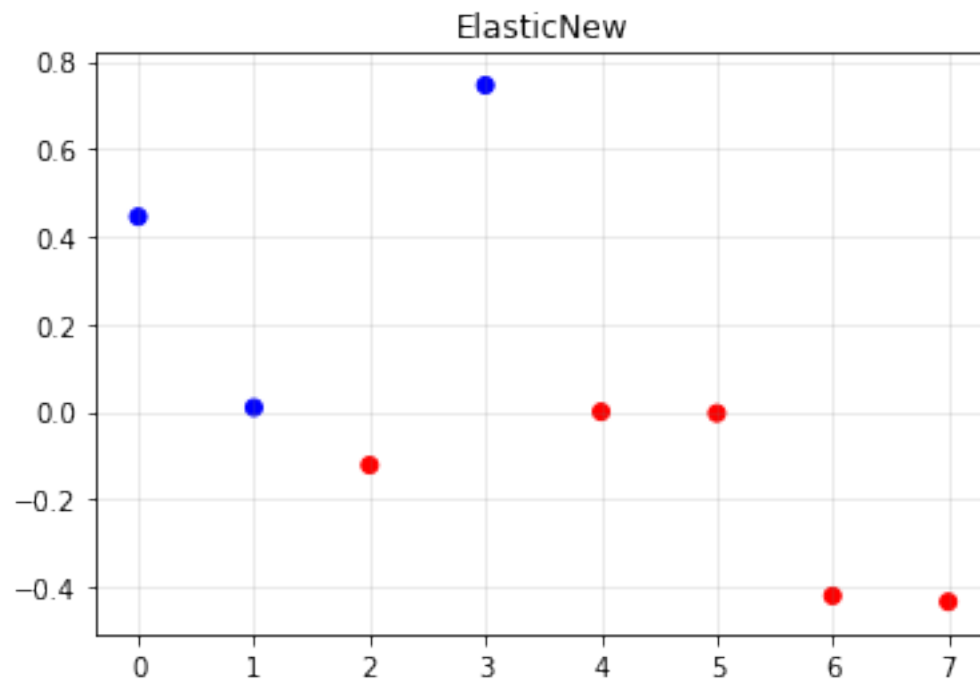
plt.scatter(range(X_train.shape[1]), lasso.coef_, c=np.sign(lasso.coef_), cmap="bwr_r")
plt.title('Lasso')
plt.show()

plt.scatter(range(X_train.shape[1]), en.coef_, c=np.sign(en.coef_), cmap="bwr_r")
plt.title('ElasticNew')
plt.show()
```









Yes, it agrees on the features which are important.

# Task2

February 7, 2018

## 1 Task2


```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import sklearn.datasets
from sklearn.svm import LinearSVC
from sklearn.neighbors import NearestCentroid
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, cross_val_score, KFold
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

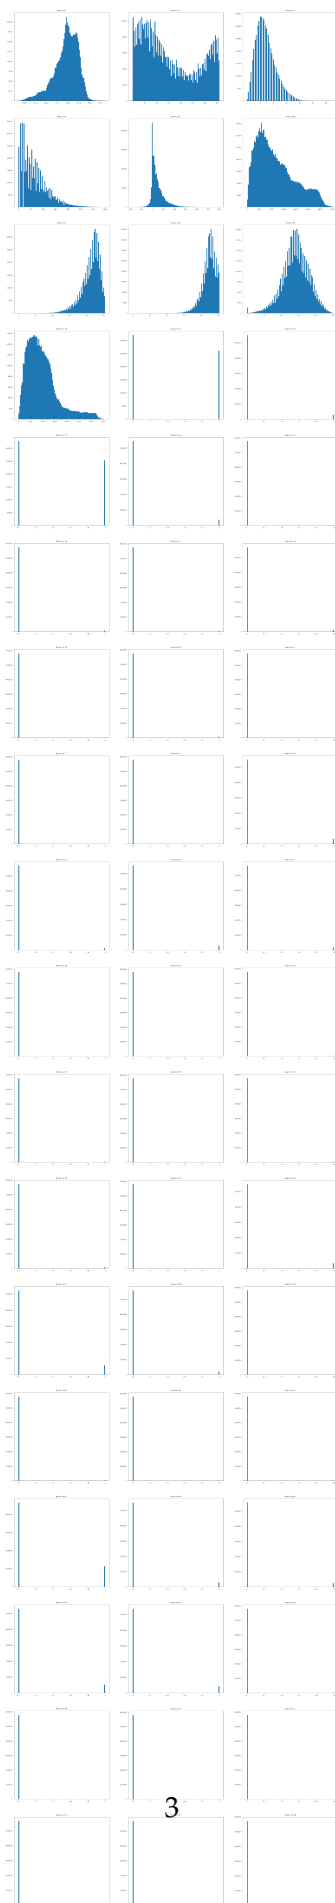
```
In [2]: covtype_dataset = sklearn.datasets.fetch_covtype()
```

### 2.1

```
In [7]: plt.hist(covtype_dataset['target'], bins=100)
plt.grid(color='gray', linestyle='-', linewidth=0.5, alpha=0.3)
plt.show()

fig, axes = plt.subplots(18, 3, figsize=(30, 180))
for i in range(0, 18):
    for j in range(0, 3):
        axes[i, j].hist(covtype_dataset['data'][:, i*3+j], bins=100)
        axes[i, j].set_title('feature{}'.format(i*3+j+1))
plt.show()
```





## 2.2

```
In [10]: X_train, X_test, y_train, y_test = train_test_split(covtype_dataset['data'], covtype_data,
    LRScore=np.mean(cross_val_score(LogisticRegression(tol=0.1, dual=False, solver='sag'), X_train, y_train)),
    LinearSVCScore=np.mean(cross_val_score(LinearSVC(tol=0.1, dual=False), X_train, y_train)),
    NearestCentroidScore=np.mean(cross_val_score(NearestCentroid(), X_train, y_train)))
```

```
print('LRScore:{}\nLinearSVCScore:{}\nNearestCentroidScore:{}\n'.format(LRScore, LinearSVCScore, NearestCentroidScore))
```

LRScore:0.6331321130115436

LinearSVCScore:0.528099269612307

NearestCentroidScore:0.19453872731809618

Scaling the Features:

```
In [11]: scaler = StandardScaler()
    scaler.fit(X_train)
    X_train_scaled = scaler.transform(X_train)
    LRScore_S=np.mean(cross_val_score(LogisticRegression(tol=0.1, dual=False, solver='sag'), X_train_scaled, y_train)),
    LinearSVCScore_S=np.mean(cross_val_score(LinearSVC(tol=0.1, dual=False), X_train_scaled, y_train)),
    NearestCentroidScore_S=np.mean(cross_val_score(NearestCentroid(), X_train_scaled, y_train)))
```

```
print('LRScore:{}\nLinearSVCScore:{}\nNearestCentroidScore:{}\n'.format(LRScore_S, LinearSVCScore_S, NearestCentroidScore_S))
```

LRScore:0.7147207572598478


LinearSVCScore:0.712627852387348

NearestCentroidScore:0.5508319986457934

After scaling, the scores get highly promoted. Scaling works.

## 2.3


```
In [24]: # LR
    param_grid = {'C': np.logspace(-3, 3, 13)}
    grid = GridSearchCV(LogisticRegression(tol=0.1, dual=False, solver='sag'), param_grid, cv=5)
    grid.fit(X_train, y_train)
    lr = grid.best_estimator_
    plt.plot(param_grid['C'], grid.cv_results_['mean_train_score'], c='blue', label='mean_train_score')
    plt.plot(param_grid['C'], grid.cv_results_['mean_test_score'], c='red', label='mean_test_score')
    plt.xlabel('LR C')
    plt.ylabel('mean cv score')
    plt.xscale('log')
    plt.legend()
    plt.show()
    print('best score:{}\nbest parameters:{}'.format(grid.best_score_, grid.best_params_))
```



best score:0.6402919962639899  
best parameters: {'C': 0.01}

In [25]: #SVM


```
param_grid = {'C': np.logspace(-5, -3, 13)}
grid = GridSearchCV(LinearSVC(tol=0.1,dual=False), param_grid,return_train_score=True)
grid.fit(X_train, y_train)
svm = grid.best_estimator_
plt.plot(param_grid['C'], grid.cv_results_['mean_train_score'],c='blue',label='mean_train_score')
plt.plot(param_grid['C'], grid.cv_results_['mean_test_score'],c='red',label='mean_test_score')
plt.xlabel('SVM C')
plt.ylabel('mean cv score')
plt.xscale('log')
plt.legend()
plt.show()
print('best score:{}\nbest parameters:{}'.format(grid.best_score_,grid.best_params_))
```



best score:0.528101542366308  
best parameters:{'C': 1e-05}

In [22]: # NC

```
param_grid = {'shrink_threshold': np.logspace(-3, 4, 10)}
grid = GridSearchCV(NearestCentroid(), param_grid, return_train_score=True)
grid.fit(X_train, y_train)
svm = grid.best_estimator_
plt.plot(param_grid['shrink_threshold'], grid.cv_results_['mean_train_score'], c='blue')
plt.plot(param_grid['shrink_threshold'], grid.cv_results_['mean_test_score'], c='red')
plt.xlabel('shrink_threshold')
plt.ylabel('mean cv score')
plt.xscale('log')
plt.legend()
plt.show()
print('best score:{}'.format(grid.best_score_))
print('best parameters:{}'.format(grid.best_params_))
```



best score:0.4137172152497137

best parameters: {'shrink\_threshold': 278.2559402207126}

According to the results, the GridSearch only works on the NearestCentroid model while it nearly has no influence on other two models.

## 2.4

Change the cross-validation strategy from 'stratified k-fold' to 'kfold' with shuffling

```
In [13]: kf=KFold(shuffle=True)
param_grid = {'C': np.logspace(-3, 3, 13)}
grid = GridSearchCV(LogisticRegression(tol=0.1,dual=False,solver='sag'), param_grid, cv=kf)
grid.fit(X_train, y_train)
lr = grid.best_estimator_

print('best score:{}\nbest parameters:{}'.format(grid.best_score_,grid.best_params_))
#SVM

param_grid = {'C': np.logspace(-5, -3, 13)}
grid = GridSearchCV(LinearSVC(tol=0.1,dual=False), param_grid, cv=kf)
grid.fit(X_train, y_train)
svm = grid.best_estimator_

print('best score:{}\nbest parameters:{}'.format(grid.best_score_,grid.best_params_))
# NC
```



```

param_grid = {'shrink_threshold': np.logspace(-3, 4, 10)}
grid = GridSearchCV(NearestCentroid(), param_grid, cv=kf)
grid.fit(X_train, y_train)
svm = grid.best_estimator_

print('best score:{}\nbest parameters:{}'.format(grid.best_score_, grid.best_params_))

best score:0.640523775756783
best parameters: {'C': 0.0031622776601683794}
best score:0.5331043076562962
best parameters: {'C': 1e-05}
best score:0.4137860606436126
best parameters: {'shrink_threshold': 278.2559402207126}

```

After changing the cross-validation strategy from 'stratified k-fold' to 'kfold' with shuffling, the parameter of Logistic Regression and linear support vector machines have changed while that of nearest centroids doesn't change at all.

Change the random seed of the shuffling

```

In [14]: kf=KFold(random_state=8)
param_grid = {'C': np.logspace(-3, 3, 13)}
grid = GridSearchCV(LogisticRegression(tol=0.1,dual=False,solver='sag'), param_grid, cv=kf)
grid.fit(X_train, y_train)
lr = grid.best_estimator_

print('best score:{}\nbest parameters:{}'.format(grid.best_score_, grid.best_params_))
#SVM
param_grid = {'C': np.logspace(-5, -3, 13)}
grid = GridSearchCV(LinearSVC(tol=0.1,dual=False), param_grid, cv=kf)
grid.fit(X_train, y_train)
svm = grid.best_estimator_

print('best score:{}\nbest parameters:{}'.format(grid.best_score_, grid.best_params_))
# NC
param_grid = {'shrink_threshold': np.logspace(-3, 4, 10)}
grid = GridSearchCV(NearestCentroid(), param_grid, cv=kf)
grid.fit(X_train, y_train)
svm = grid.best_estimator_

print('best score:{}\nbest parameters:{}'.format(grid.best_score_, grid.best_params_))

best score:0.6431031831815293
best parameters: {'C': 3.1622776601683795}
best score:0.5316493749985657
best parameters: {'C': 0.001}
best score:0.41334086042973295
best parameters: {'shrink_threshold': 278.2559402207126}

```

After changing random seed of the shuffling, the parameter of Logistic Regression and inear support vector machines have changed while that of nearest centroids doesn't change at all.  
Change the random state of the split into training and test data

```
In [15]: kf=KFold(n_splits=2)
```

```
param_grid = {'C': np.logspace(-3, 3, 13)}
grid = GridSearchCV(LogisticRegression(tol=0.1,dual=False,solver='sag'), param_grid, cv=kf)
grid.fit(X_train, y_train)
lr = grid.best_estimator_

print('best score:{}\nbest parameters:{}'.format(grid.best_score_,grid.best_params_))
#SVM
param_grid = {'C': np.logspace(-5, -3, 13)}
grid = GridSearchCV(LinearSVC(tol=0.1,dual=False), param_grid, cv=kf)
grid.fit(X_train, y_train)
svm = grid.best_estimator_

print('best score:{}\nbest parameters:{}'.format(grid.best_score_,grid.best_params_))
# NC
param_grid = {'shrink_threshold': np.logspace(-3, 4, 10)}
grid = GridSearchCV(NearestCentroid(), param_grid, cv=kf)
grid.fit(X_train, y_train)
nc = grid.best_estimator_

print('best score:{}\nbest parameters:{}'.format(grid.best_score_,grid.best_params_))
```

```
best score:0.6352112061942495
best parameters: {'C': 0.01}
best score:0.4979013629093145
best parameters: {'C': 1e-05}
best score:0.47876922794480437
best parameters: {'shrink_threshold': 278.2559402207126}
```

After changing random state of the split into training and test data, the parameter of Logistic Regression and inear support vector machines have changed while that of nearest centroids doesn't change at all.

2.5

Choose the model with the best performance, the model after scaling the data with Standard-Scaler.

```
In [18]: scaler = StandardScaler()
```

```
scaler.fit(X_train)
```

```
X_train_scaled = scaler.transform(X_train)
```

```
lr = LogisticRegression(tol=0.1,dual=False,solver='sag').fit(X_train_scaled, y_train)
```

```
for i in range(0,7):
```

```
    plt.scatter(range(X_train_scaled.shape[1]), lr.coef_[i],c=np.sign(lr.coef_[i]), cmap="bw")
```

```

plt.title('Logistic Regression set{}'.format(i))
plt.show()
svm = LinearSVC(tol=0.1,dual=False).fit(X_train, y_train)
for i in range(0,7):
    plt.scatter(range(X_train.shape[1]),svm.coef_[i], c=np.sign(svm.coef_[i]), cmap="")
    plt.ylim(-0.005,0.005)
    plt.title('SVM set{}'.format(i))
    plt.show()

```

