

Task1

February 7, 2018

1 Task1

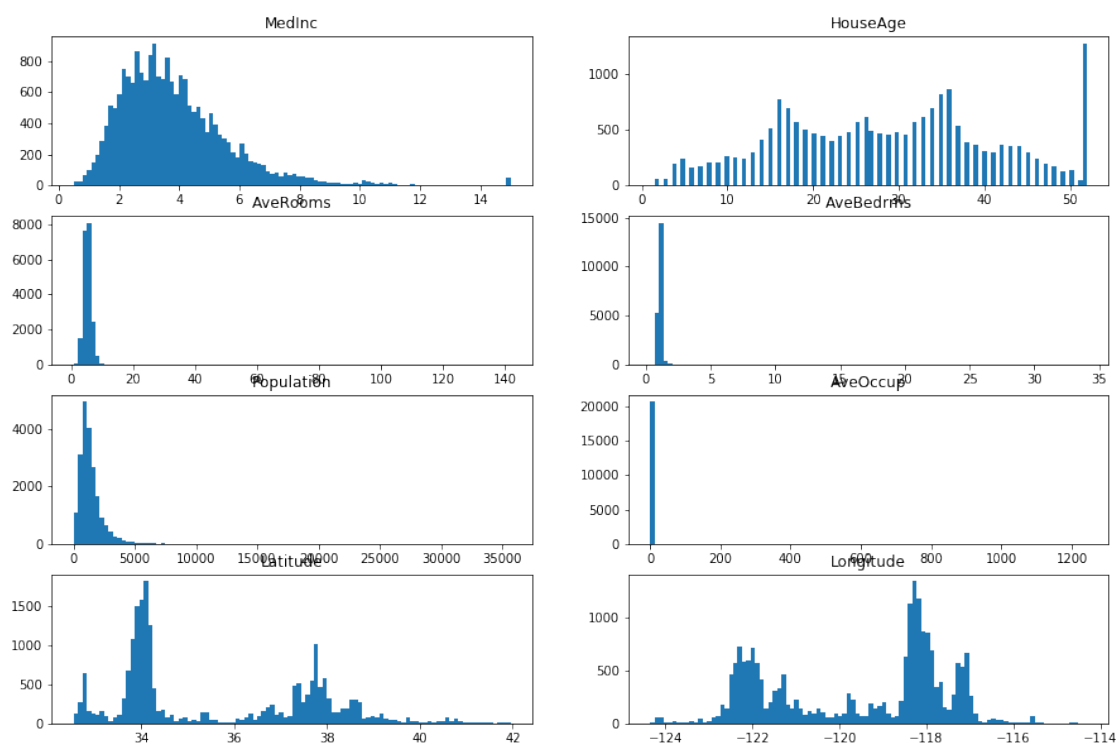

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import sklearn.datasets
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

1.1

```
In [18]: california_dataset = sklearn.datasets.fetch_california_housing()

plt.hist(california_dataset['target'], bins=100)
plt.grid(color='gray', linestyle='-', linewidth=0.5, alpha=0.3)
plt.title('target')
plt.show()

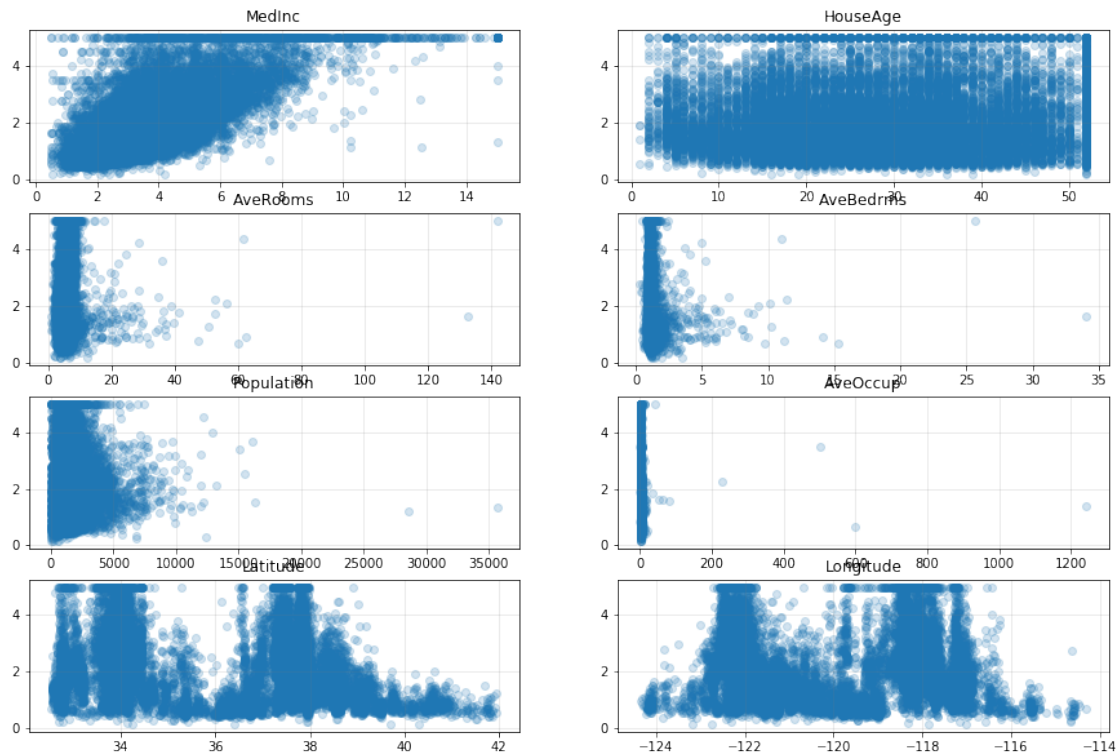
fig, axes = plt.subplots(4, 2, figsize=(15, 10))
for i in range(0, 4):
    for j in range(0, 2):
        axes[i, j].hist(california_dataset['data'][:, i*2+j], bins=100)
        axes[i, j].set_title(california_dataset['feature_names'][i*2+j])
plt.show()
```



From the above eight feature plots, 'AveRooms', 'AveBedrms', 'Population' and 'AveOccup' have inappropriate x scale, which means there are some outliers on these four features. We may need to remove these outliers, than re-scale the x axis.

1.2

```
In [4]: fig, axes = plt.subplots(4,2,figsize=(15,10))
        for i in range(0,4):
            for j in range(0,2):
                axes[i,j].scatter(california_dataset['data'][:,i*2+j],california_dataset['target'])
                axes[i,j].set_title(california_dataset['feature_names'][i*2+j])
        plt.show()
```



1.3

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(california_dataset['data'],california_dataset['target'],
                    LRScore=np.mean(cross_val_score(LinearRegression(), X_train, y_train))
                    RidgeScore=np.mean(cross_val_score(Ridge(), X_train, y_train))
                    LassoScore=np.mean(cross_val_score(Lasso(), X_train, y_train))
                    ElasticNetScore=np.mean(cross_val_score(ElasticNet(), X_train, y_train))
                    print('LRScore:{}\nRidgeScore:{}\nLassoScore:{}\nElasticNetScore:{}'.format(LRScore,RidgeScore,LassoScore,ElasticNetScore))
```

LRScore:0.6102422922538064

RidgeScore:0.6102437749800811

```
LassoScore:0.2819718569982819
ElasticNetScore:0.4229514321435757
```

Scaling the features

```
In [13]: scaler = StandardScaler()
        scaler.fit(X_train)
        X_train_scaled = scaler.transform(X_train)
        LRScore_S=np.mean(cross_val_score(LinearRegression(), X_train_scaled, y_train))
        RidgeScore_S=np.mean(cross_val_score(Ridge(), X_train_scaled, y_train))
        LassoScore_S=np.mean(cross_val_score(Lasso(), X_train_scaled, y_train))
        ElasticNetScore_S=np.mean(cross_val_score(ElasticNet(), X_train_scaled, y_train))
        print('LRScore:{}\nRidgeScore:{}\nLassoScore:{}\nElasticNetScore:{}'.format(LRScore_S,
        RidgeScore_S, LassoScore_S, ElasticNetScore_S))

LRScore:0.6102422922538069
RidgeScore:0.6102441875598729
LassoScore:-0.00015307833209199373
ElasticNetScore:0.20448945641200397
```

According to the results, Scaling doesn't help the OLR and Ridge. Besides, it even has a negative influence on the Lasso and ElasticNet.

1.4

```
In [20]: # Ridge
        param_grid = {'alpha': np.logspace(-3, 3, 13)}
        grid = GridSearchCV(Ridge(), param_grid, return_train_score=True)
        grid.fit(X_train, y_train)
        plt.plot(param_grid['alpha'], grid.cv_results_['mean_train_score'], c='blue', label='mean_train_score')
        plt.plot(param_grid['alpha'], grid.cv_results_['mean_test_score'], c='red', label='mean_test_score')
        plt.xlabel('alpha')
        plt.ylabel('mean cv score')
        plt.xscale('log')
        plt.legend()
        plt.show()
        ridge = grid.best_estimator_
        print('best score:{}\nbest parameters:{}'.format(grid.best_score_, grid.best_params_))

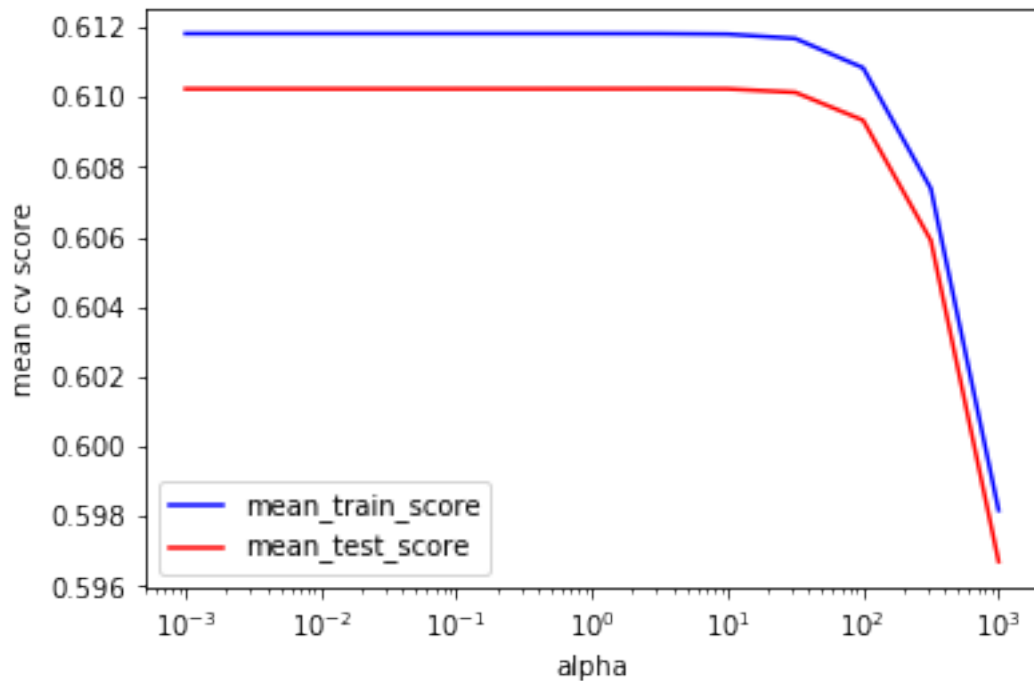
#Lasso
        param_grid = {'alpha': np.logspace(-3, 0, 13)}
        grid = GridSearchCV(Lasso(), param_grid, return_train_score=True)
        grid.fit(X_train, y_train)
        plt.plot(param_grid['alpha'], grid.cv_results_['mean_train_score'], c='blue', label='mean_train_score')
        plt.plot(param_grid['alpha'], grid.cv_results_['mean_test_score'], c='red', label='mean_test_score')
        plt.xlabel('alpha')
        plt.ylabel('mean cv score')
```

```

plt.xscale('log')
plt.legend()
plt.show()
lasso = grid.best_estimator_
print('best score:{}'.format(grid.best_score_))

# ElasticNet
param_grid = {'alpha': np.logspace(-3, 2, 10), 'l1_ratio': [0.01, .1, .5, .9, .98, 1]}
grid = GridSearchCV(ElasticNet(), param_grid, return_train_score=True)
grid.fit(X_train, y_train)
res = pd.pivot_table(pd.DataFrame(grid.cv_results_), values='mean_test_score', index='alpha')
plt.imshow(res, extent=[0, 1, 0, 100], aspect="auto")
plt.colorbar()
plt.show()
en = grid.best_estimator_
print('best score:{}'.format(grid.best_score_))

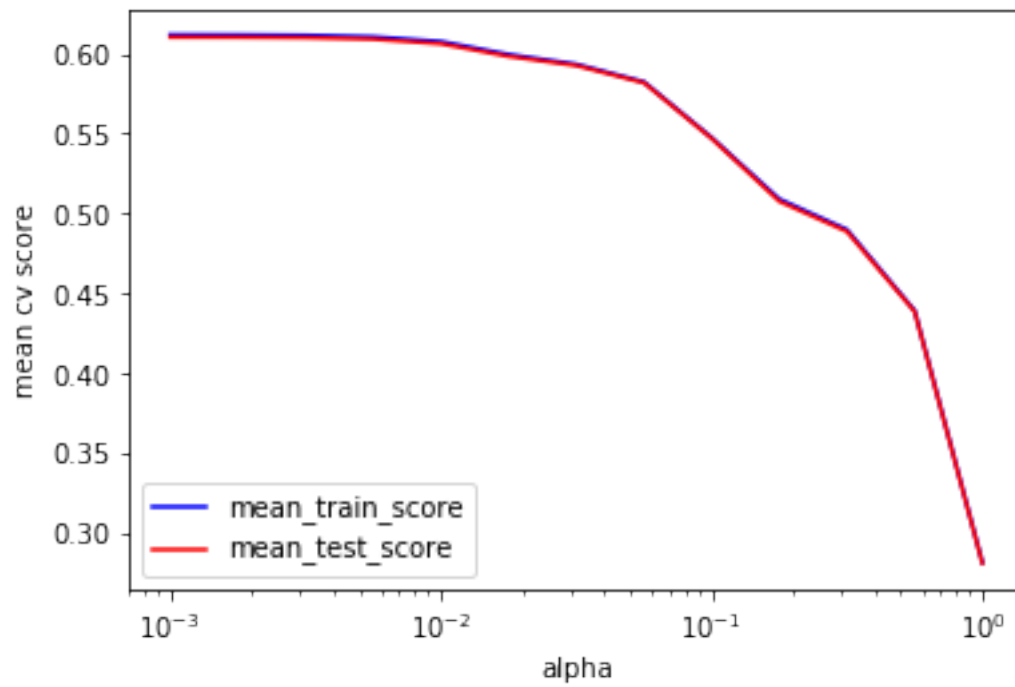
```




```

best score:0.6102457318944084
best parameters: {'alpha': 3.1622776601683795}

```



best score:0.6102299335181556
best parameters: {'alpha': 0.001}



```
best score:0.6102407928395485
best parameters: {'alpha': 0.001, 'l1_ratio': 0.01}
```

According to the results, GridSearch only can improve Lasso and ElasticNet.

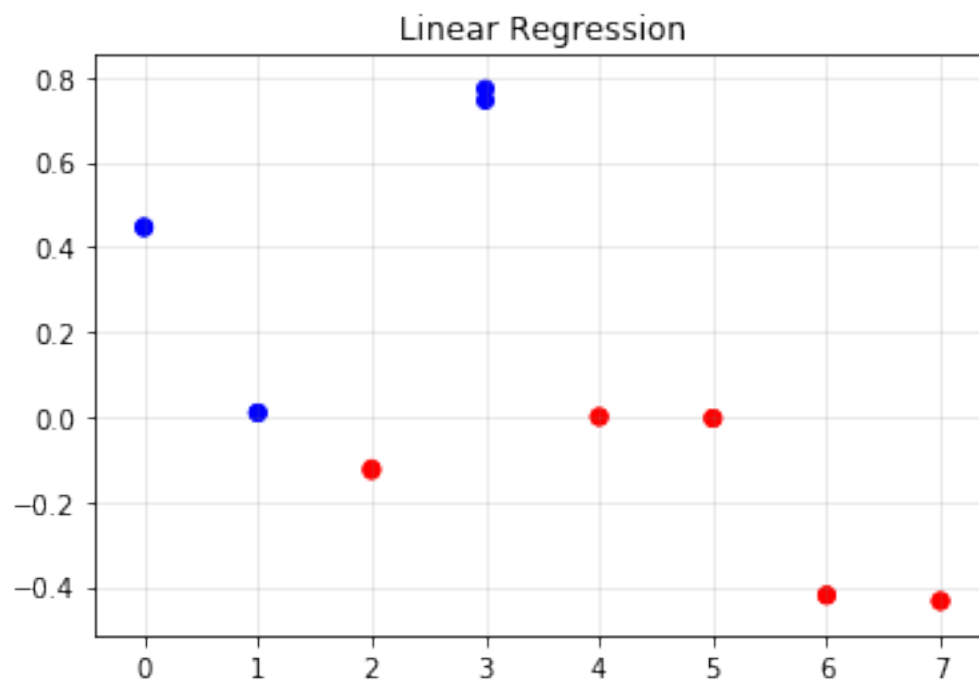
1.5

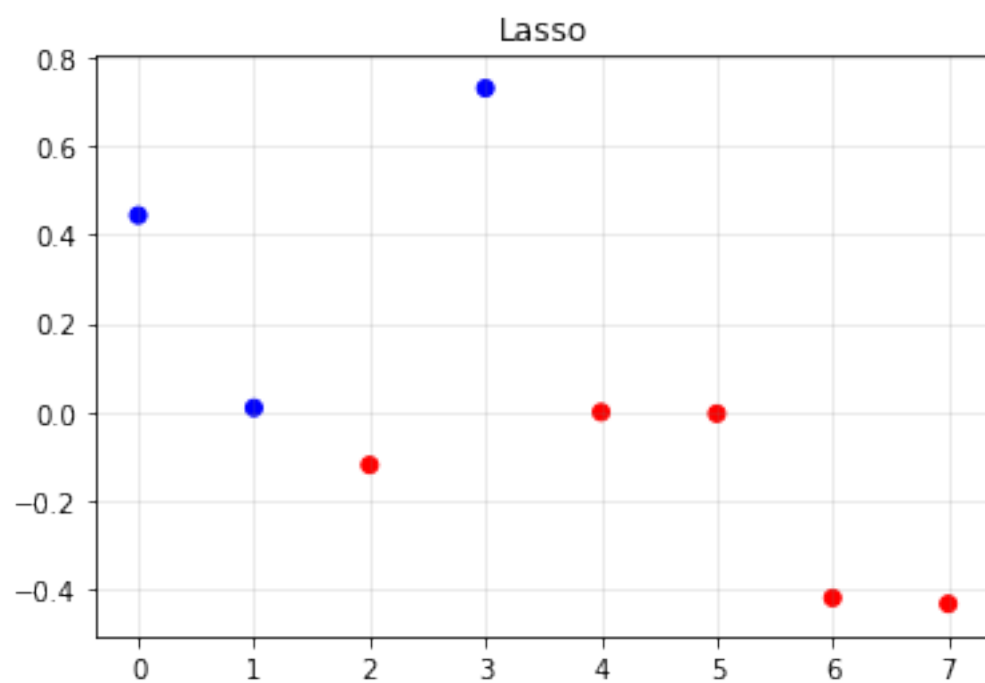
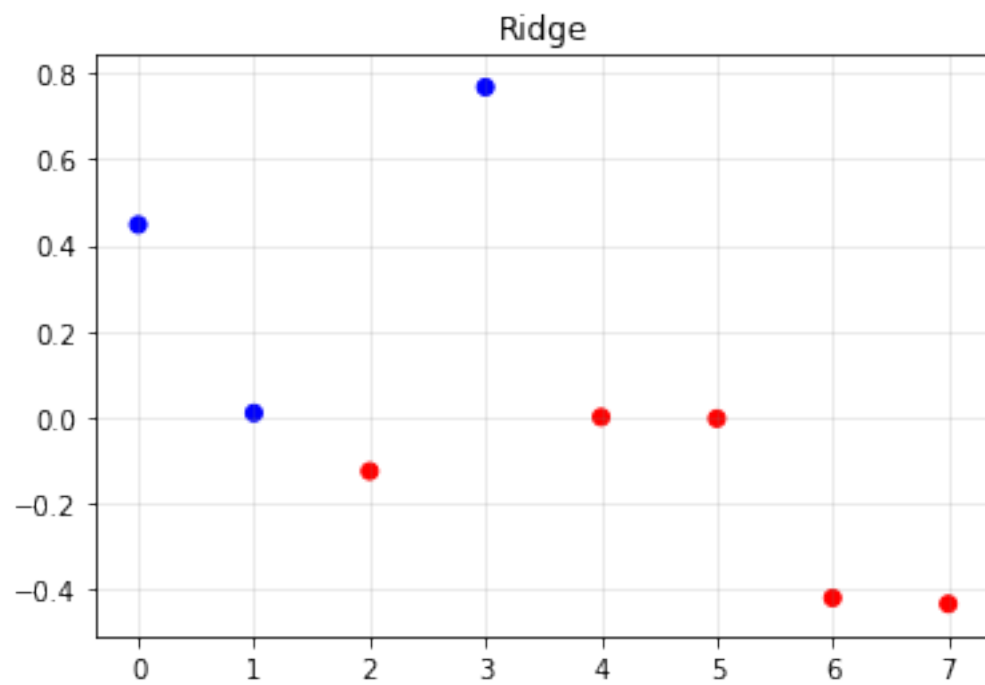
```
In [17]: lr = LinearRegression().fit(X_train, y_train)
plt.scatter(range(X_train.shape[1]), lr.coef_, c=np.sign(lr.coef_), cmap="bwr_r")
plt.title('Linear Regression')
plt.show()


plt.scatter(range(X_train.shape[1]), ridge.coef_, c=np.sign(ridge.coef_), cmap="bwr_r")
plt.title('Ridge')
plt.show()

plt.scatter(range(X_train.shape[1]), lasso.coef_, c=np.sign(lasso.coef_), cmap="bwr_r")
plt.title('Lasso')
plt.show()

plt.scatter(range(X_train.shape[1]), en.coef_, c=np.sign(en.coef_), cmap="bwr_r")
plt.title('ElasticNew')
plt.show()
```







Yes, it agrees on the features which are important.