

AML_HW5 2

April 16, 2018

```
#  
AML HW 5  
####  
April 16, 2018  
Jingyu Ren UNI: jr3738  
Chi Zhang UNI: cz2481
```

```
In [1]: import pandas as pd  
import numpy as np  
from sklearn.feature_extraction.text import CountVectorizer  
from sklearn.linear_model import LogisticRegressionCV, LogisticRegression  
from sklearn.pipeline import make_pipeline  
from sklearn.model_selection import cross_val_score  
from scipy.sparse import hstack  
import matplotlib.pyplot as plt  
from sklearn.model_selection import GridSearchCV  
from sklearn.feature_extraction.text import TfidfVectorizer  
from sklearn.preprocessing import StandardScaler  
from sklearn.preprocessing import Normalizer
```

1 Task 1

```
In [2]: train = pd.DataFrame.from_csv('/Users/jingyu/Desktop/hw5_data_train.csv', index_col=None)  
test = pd.DataFrame.from_csv('/Users/jingyu/Desktop/hw5_data_test.csv', index_col=None)  
  
y_train = train['Recommended']  
y_test = test['Recommended']  
  
#1) Use the title only  
vect = CountVectorizer()  
title_train = vect.fit_transform(train['Title'])  
name1 = vect.get_feature_names()  
  
#2) Use the review body only  
vect = CountVectorizer()  
review_train = vect.fit_transform(train['Review'])  
name2 = vect.get_feature_names()
```

```

#3) Concatenate the title and review to a single text and analyze that (discarding the
vect = CountVectorizer()
titleReview_train = vect.fit_transform(train['Title'].map(str) + ' ' + train['Review'])
name3 = vect.get_feature_names()
titleReview_test = test['Title'].map(str) + ' ' + test['Review']

#4) Vectorizing title and review individually and concatenating the vector representat
vect1 = CountVectorizer()
vect2 = CountVectorizer()
title_review_train = hstack((vect1.fit_transform(train['Title']), vect2.fit_transform(t
name4 = name1 + name2

```

```

/Users/jingyu/anaconda/envs/python3/lib/python3.5/site-packages/ipykernel_launcher.py:1: Future
"""Entry point for launching an IPython kernel.
/Users/jingyu/anaconda/envs/python3/lib/python3.5/site-packages/ipykernel_launcher.py:2: Future

```

```

In [176]: pipe1 = make_pipeline(CountVectorizer(), LogisticRegression())
          print('#1 Title Only cross validation score:\n{}'.format(np.mean(cross_val_score(Log

          pipe2 = make_pipeline(CountVectorizer(), LogisticRegression())
          print('#2 Review Only cross validation score:\n{}'.format(np.mean(cross_val_score(Log

          pipe3 = make_pipeline(CountVectorizer(), LogisticRegression())
          print('#3 Concatenate Title and Review cross validation score:\n{}'.format(np.mean(c

          print('#4 Concatenate Title and Review Vectors cross validation score:\n{}'.format(np

#1 Title Only cross validation score:
0.920447284709158
#2 Review Only cross validation score:
0.9119736515693934
#3 Concatenate Title and Review cross validation score:
0.933310639278875
#4 Concatenate Title and Review Vectors cross validation score:
0.9383480297268004

```

```



In [8]: def plot_coefficients(classifier, feature_names, top_features=20):
        coef = classifier.coef_.ravel()
        top_positive_coefficients = np.argsort(coef)[-top_features:]
        top_negative_coefficients = np.argsort(coef)[:top_features]
        top_coefficients = np.hstack([top_negative_coefficients, top_positive_coefficients])
        plt.figure(figsize=(15, 5))
        colors = ['red' if c < 0 else 'blue' for c in coef[top_coefficients]]
        plt.bar(np.arange(2 * top_features), coef[top_coefficients], color=colors)
        feature_names = np.array(feature_names)



```

```
plt.xticks(np.arange(1, 1 + 2 * top_features), feature_names[top_coefficients], rotation=90)
plt.show()
```

```
In [182]: xList=[title_train, review_train, titleReview_train, title_review_train]
names=[name1,name2,name3,name4]
```

```
for i in range(4):
    lr = LogisticRegression()
    lr.fit(xList[i], y_train)
    plot_coefficients(lr, names[i])
    i = i + 1
```





1.1 GridSearch to tune the regularization parameter

```
In [301]: y_train = train['Recommended']
          y_test = test['Recommended']
```

#1) Use the title only

```
title_train = train['Title']
title_test = test['Title']
```

#2) Use the review body only

```
review_train = train['Review']
review_test = test['Review']
```

#3) Concatenate the title and review to a single text and analyze that (discarding t

```

titleReview_train = train['Title'].map(str) + ' ' + train['Review']
titleReview_test = test['Title'].map(str) + ' ' + test['Review']

#4) Vectorizing title and review individually and concatenating the vector represent
vect1 = CountVectorizer()
vect2 = CountVectorizer()
title_review_train = hstack((vect1.fit_transform(train['Title']), vect2.fit_transform
title_review_test = hstack((vect1.transform(test['Title']), vect2.transform(test['Rev

In [303]: grid = {'logisticregression__C': [0.01, 0.1, 1, 10, 100]}
pipe = make_pipeline(CountVectorizer(), LogisticRegression())
gs = GridSearchCV(pipe, grid, scoring='roc_auc', cv=5)
gs.fit(title_train, y_train)
print ('#1 best_parameter_', gs.best_params_)
print ('#1 best_cv_score_', gs.best_score_)
print ('#1 predict_score_', gs.score(title_test, y_test))

grid = {'logisticregression__C': [0.01, 0.1, 1, 10, 100]}
pipe = make_pipeline(CountVectorizer(), LogisticRegression())
gs = GridSearchCV(pipe, grid, scoring='roc_auc', cv=5)
gs.fit(review_train, y_train)
print ('#2 best_parameter_', gs.best_params_)
print ('#2 best_cv_score_', gs.best_score_)
print ('#2 predict_score_', gs.score(review_test, y_test))

grid = {'logisticregression__C': [0.01, 0.1, 1, 10, 100]}
pipe = make_pipeline(CountVectorizer(), LogisticRegression())
gs = GridSearchCV(pipe, grid, scoring='roc_auc', cv=5)
gs.fit(titleReview_train, y_train)
print ('#3 best_parameter_', gs.best_params_)
print ('#3 best_cv_score_', gs.best_score_)
print ('#3 predict_score_', gs.score(titleReview_test, y_test))

grid = {'C': [0.01, 0.1, 1, 10, 100]}
lr = LogisticRegression()
gs = GridSearchCV(lr, grid, scoring='roc_auc', cv=5)
gs.fit(title_review_train, y_train)
print ('#4 best_parameter_', gs.best_params_)
print ('#4 best_score_', gs.best_score_)
print ('#4 predict_score_', gs.score(title_review_test, y_test))



#1 best_parameter_: {'logisticregression__C': 1}
#1 best_cv_score_: 0.9204469952584714
#1 predict_score_: 0.9214984071094153
#2 best_parameter_: {'logisticregression__C': 0.1}
#2 best_cv_score_: 0.923761847972498
#2 predict_score_: 0.9221422856845667
#3 best_parameter_: {'logisticregression__C': 0.1}



```

```
#3 best_cv_score_: 0.942410344658611
#3 predict_score_: 0.9388955774186873
#4 best_parameter_: {'C': 0.1}
#4 best_score_: 0.9436032417207872
#4 predict_score_: 0.9431204167740951
```

```
In [184]: xList=[title_train, review_train, titleReview_train, title_review_train]
          names=[name1,name2,name3,name4]
          params=[1,0.1,0.1,0.1]
```

```
for i in range(4):
    lr = LogisticRegression(C=params[i])
    lr.fit(xList[i], y_train)
    plot_coefficients(lr, names[i])
    i = i + 1
```





1.1.1 Based on the above results, the fourth way is best.

2 Task 2

2.1 2.1 TfidfVectorizer

```
In [5]: titleReview_train = train['Title'].map(str) + ' ' + train['Review']
        titleReview_val = test['Title'].map(str) + ' ' + test['Review']
```

```
In [15]: param_grid_log = { 'logisticregression__C': [0.01,0.1,1,10]}
        pipe_log_vect = make_pipeline(CountVectorizer(), LogisticRegression(), memory="cache_")
        grid_log_vect = GridSearchCV(pipe_log_vect, param_grid_log, scoring='roc_auc',cv=5)
        grid_log_vect.fit(titleReview_train, y_train)
```

```
Out[15]: GridSearchCV(cv=5, error_score='raise',
                    estimator=Pipeline(memory='cache_folder',
```

```

steps=[('countvectorizer', CountVectorizer(analyzer='word', binary=False, decode
dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
lowercase=True, max_df=1.0, max_features=None, min_df=1,
ngram_range=(1, 1), preprocessor=None, stop_words=None,
...ty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False))),
fit_params=None, iid=True, n_jobs=1,
param_grid={'logisticregression__C': [0.01, 0.1, 1, 10]},
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring='roc_auc', verbose=0)

```

```

In [16]: param_tfidf = { 'logisticregression__C': [0.01, 0.1, 1, 10, 100]}
pipe_tfidf = make_pipeline(TfidfVectorizer(), LogisticRegression(), memory="cache_fold
grid_tfidf = GridSearchCV(pipe_tfidf, param_tfidf, scoring='roc_auc', cv=5)
grid_tfidf.fit(titleReview_train, y_train)
print ('grid_log_vect best_score_', grid_log_vect.best_score_)
print ('grid_log_vect predict_score_', grid_log_vect.score(titleReview_test, y_test))
print ('grid_tfidf best_parameter_', grid_tfidf.best_params_)
print ('grid_tfidf best_score_', grid_tfidf.best_score_)
print ('grid_tfidf predict_score_', grid_tfidf.score(titleReview_test, y_test))

```

```


grid_log_vect best_score_: 0.942410344658611
grid_log_vect predict_score_: 0.9388955774186873
grid_tfidf best_parameter_: {'logisticregression__C': 1}
grid_tfidf best_score_: 0.9498551621199149
grid_tfidf predict_score_: 0.9454635538445706


```

```

In [19]: plot_coefficients(grid_log_vect.best_estimator_.named_steps['logisticregression'], gr
plot_coefficients(grid_tfidf.best_estimator_.named_steps['logisticregression'], grid_t

```







Compared with part 1.3, using `tfidfvectorizer`, the cv score on training set has been improved from 0.9424 to 0.9499. The predict score has benn improve from 0.9389 to 0.9456. Also the important features with related coefficients changed. For example, the new words has high weight such as love, fits and so on.

2.2 Normalizer with CountVectorizer

```
In [17]: param_norm = { 'logisticregression__C': [0.01, 0.1, 1, 10, 100]}
pipe_norm = make_pipeline(CountVectorizer(), Normalizer(), LogisticRegression(), memo
grid_norm = GridSearchCV(pipe_norm, param_norm, scoring='roc_auc', cv=5)
grid_norm.fit(titleReview_train, y_train)
print ('grid_tfidf best_score_', grid_tfidf.best_score_)
print ('grid_tfidf predict_score_', grid_tfidf.score(titleReview_test, y_test))
print ('grid_norm best_parameter_', grid_norm.best_params_)
print ('grid_norm best_score_', grid_norm.best_score_)
print ('grid_norm predict_score_', grid_norm.score(titleReview_test, y_test))
```

```
grid_tfidf best_score_: 0.9498551621199149
grid_tfidf predict_score_: 0.9454635538445706
grid_norm best_parameter_: {'logisticregression__C': 10}
grid_norm best_score_: 0.9476031785092168
grid_norm predict_score_: 0.9441320876438387
```

```
In [20]: plot_coefficients(grid_tfidf.best_estimator_.named_steps['logisticregression'], grid_t
plot_coefficients(grid_norm.best_estimator_.named_steps['logisticregression'], grid_n
```



the CV score and predict score didn't improve. But the important feature coef changed.

2.3 2.3 Stop-word

```
In [21]: param_stop = { 'logisticregression_C': [0.01, 0.1, 1, 10, 100]}
pipe_stop = make_pipeline(TfidfVectorizer(stop_words='english'), LogisticRegression())
grid_stop = GridSearchCV(pipe_stop, param_stop, scoring='roc_auc', cv=5)
grid_stop.fit(titleReview_train, y_train)
print ('grid_tfidf best_score:', grid_tfidf.best_score_)
print ('grid_tfidf predict_score:', grid_tfidf.score(titleReview_test, y_test))
print ('grid_stop best_parameter:', grid_stop.best_params_)
print ('grid_stop best_score:', grid_stop.best_score_)
print ('grid_stop predict_score:', grid_stop.score(titleReview_test, y_test))
```

```



grid_tfidf best_score_: 0.9498551621199149
grid_tfidf predict_score_: 0.9454635538445706
grid_stop best_parameter_: {'logisticregression__C': 1}
grid_stop best_score_: 0.9434995623247132
grid_stop predict_score_: 0.9397407112735047

```

```

In [22]: plot_coefficients(grid_tfidf.best_estimator_.named_steps['logisticregression'], grid_tfidf.best_estimator_.named_steps['logisticregression'], grid_stop.best_estimator_.named_steps['logisticregression'], grid_stop.best_estimator_.named_steps['logisticregression'])

```



the CV score and predict score didn't improve. But the important feature coef didn't change.

2.4 min_df

```

In [23]: param4min = { 'logisticregression__C': [0.01, 0.1, 1, 10, 100]}
pipe4min = make_pipeline(TfidfVectorizer(min_df=4), LogisticRegression(), memory="cache")

```

```

grid4min = GridSearchCV(pipe4min, param4min, scoring='roc_auc', cv=5)
grid4min.fit(titleReview_train, y_train)
param3min = { 'logisticregression__C': [0.01, 0.1, 1, 10, 100]}
pipe3min = make_pipeline(TfidfVectorizer(min_df=3), LogisticRegression(), memory="cache")
grid3min = GridSearchCV(pipe3min, param3min, scoring='roc_auc', cv=5)
grid3min.fit(titleReview_train, y_train)
param2min = { 'logisticregression__C': [0.01, 0.1, 1, 10, 100]}
pipe2min = make_pipeline(TfidfVectorizer(min_df=2), LogisticRegression(), memory="cache")
grid2min = GridSearchCV(pipe2min, param2min, scoring='roc_auc', cv=5)
grid2min.fit(titleReview_train, y_train)
print ('grid_tfidf best_score_', grid_tfidf.best_score_)
print ('grid_tfidf predict_score_', grid_tfidf.score(titleReview_test, y_test))
print ('grid4min best_parameter_', grid4min.best_params_)
print ('grid4min best_score_', grid4min.best_score_)
print ('grid4min predict_score_', grid4min.score(titleReview_test, y_test))
print ('grid3min best_parameter_', grid3min.best_params_)
print ('grid3min best_score_', grid3min.best_score_)
print ('grid3min predict_score_', grid3min.score(titleReview_test, y_test))
print ('grid2min best_parameter_', grid2min.best_params_)
print ('grid2min best_score_', grid2min.best_score_)
print ('grid2min predict_score_', grid2min.score(titleReview_test, y_test))

```

```




grid_tfidf best_score_: 0.9498551621199149
grid_tfidf predict_score_: 0.9454635538445706
grid4min best_parameter_: {'logisticregression__C': 1}
grid4min best_score_: 0.9500145082145014
grid4min predict_score_: 0.9454760026247561
grid3min best_parameter_: {'logisticregression__C': 1}
grid3min best_score_: 0.950026395977515
grid3min predict_score_: 0.9455927445189405
grid2min best_parameter_: {'logisticregression__C': 1}
grid2min best_score_: 0.9500371582729809
grid2min predict_score_: 0.9455227547103419


```

```

In [24]: plot_coefficients(grid_tfidf.best_estimator_.named_steps['logisticregression'], grid_tfidf.coef_)
plot_coefficients(grid4min.best_estimator_.named_steps['logisticregression'], grid4min.coef_)
plot_coefficients(grid3min.best_estimator_.named_steps['logisticregression'], grid3min.coef_)
plot_coefficients(grid2min.best_estimator_.named_steps['logisticregression'], grid2min.coef_)

```





with `min_df`, the CV score and predict score slightly improved. But the important feature coef didn't change.

3 Task 3.1

```
In [268]: X_train = train['Title'].map(str) + ' ' + train['Review']
```

```
In [271]: pipe1 = make_pipeline(TfidfVectorizer(ngram_range=(1, 1)), LogisticRegressionCV(),
print(' unigrams score: {}'.format(np.mean(cross_val_score(pipe1, X_train, y_train,


tfidf = TfidfVectorizer(ngram_range=(1, 4))
titleReview_train_tfidf = tfidf.fit_transform(train['Title'].map(str) + ' ' + train[
name = tfidf.get_feature_names()
pipe3 = make_pipeline(TfidfVectorizer(ngram_range=(1, 4)), LogisticRegressionCV(),
print('(1, 4)grams score: {}'.format(np.mean(cross_val_score(pipe3, X_train, y_train
```

```
unigrams score: 0.9467156329833519
(1, 4)grams score: 0.9563276559749114
```

Thus, n-grams(here 4-grams) of varying length will give a best performance.

```
In [207]: tfidf = TfidfVectorizer(ngram_range=(1, 4))
titleReview_train_tfidf = tfidf.fit_transform(train['Title'].map(str) + ' ' + train[
name = tfidf.get_feature_names()

lr = LogisticRegression(C=0.1)
lr.fit(titleReview_train_tfidf, y_train)
plot_coefficients(lr, name)
```



3.0.1 Draw only non-unigrams

```
In [22]: def plot_higher_coefficients(classifier, feature_names, top_features=20):
    coef = classifier.coef_.ravel()
    ngrams_name=[]
    ngrams_coef=[]
    for name in feature_names:
        if len(name.split())>1:
            ngrams_name.append(True)
        else:
            ngrams_name.append(False)
    for i in range(len(feature_names)):
        if ngrams_name[i]==True:
            ngrams_coef.append(coef[i])
        else:
            ngrams_coef.append(0)

    ngrams_coef = np.asarray(ngrams_coef)

    top_positive_coefficients = np.argsort(ngrams_coef)[-top_features:]
    top_negative_coefficients = np.argsort(ngrams_coef)[:top_features]
    top_coefficients = np.hstack([top_negative_coefficients, top_positive_coefficients])
    plt.figure(figsize=(15, 5))
    colors = ['red' if c < 0 else 'blue' for c in coef[top_coefficients]]
    plt.bar(np.arange(2 * top_features), coef[top_coefficients], color=colors)
    feature_names = np.array(feature_names)
    plt.xticks(np.arange(1, 1 + 2 * top_features), feature_names[top_coefficients], rotation=45)
    plt.show()


In [242]: tfidf = TfidfVectorizer(ngram_range=(1, 4))
    titleReview_train_tfidf = tfidf.fit_transform(train['Title'].map(str) + ' ' + train['Review'])
```

```

name = tfidf.get_feature_names()

lr = LogisticRegression(C=0.1)
lr.fit(titleReview_train_tfidf, y_train)
plot_higher_coefficients(lr, name)

```



4 Task 3.2


```

In [20]: tfidf = TfidfVectorizer(ngram_range=(1, 3), analyzer="char_wb")
titleReview_train_tfidf = tfidf.fit_transform(train['Title'].map(str) + ' ' + train['Review'])
titleReview_test_tfidf = tfidf.transform(test['Title'].map(str) + ' ' + test['Review'])
name = tfidf.get_feature_names()

In [25]: lr = LogisticRegression(C=0.1)
lr.fit(titleReview_train_tfidf, y_train)
print('Using character n-grams test score: {}'.format(lr.score(titleReview_test_tfidf, y_test)))
plot_coefficients(lr, name)

```

Using character n-grams test score: 0.8172196214125789



From the plot, we can tell that:

1. reviews or titles that have '!' in content, will lead to a recommendation.
2. reviews or titles that have 'was'-like words (such as 'wa' and 'was') in content, will not lead to a recommendation.
3. reviews or titles that have 'no'-like words (such as 'no' and 'not') in content, will not lead to a recommendation.
4. character n-grams's performance is worse than same n-grams of words.

5 Task 3.3

5.0.1 impact of min_df with n-grams

```
In [26]: X_train = train['Title'].map(str) + ' ' + train['Review']
         X_test = test['Title'].map(str) + ' ' + test['Review']

In [29]: pipe = make_pipeline(TfidfVectorizer(ngram_range=(1, 2), min_df=2), LogisticRegression)
         tfidf = TfidfVectorizer(ngram_range=(1, 2), min_df=2)
         tfidf.fit(X_train,y_train)
         print('(1, 2), min_df=2: {}'.format(len(tfidf.vocabulary_)))
         print('(1, 2), min_df=2 score: {}'.format(np.mean(cross_val_score(pipe, X_train, y_train))))
         print('(1, 2), min_df=2 test score: {}'.format(pipe.fit(X_train,y_train).score(X_test)))

         pipe = make_pipeline(TfidfVectorizer(ngram_range=(1, 2), min_df=4), LogisticRegression)
         tfidf = TfidfVectorizer(ngram_range=(1, 2), min_df=4)
         tfidf.fit(X_train,y_train)
         print("(1, 2), min_df=4: {}".format(len(tfidf.vocabulary_)))
         print('(1, 2), min_df=4 score: {}'.format(np.mean(cross_val_score(pipe, X_train, y_train))))
         print('(1, 2), min_df=4 test score: {}'.format(pipe.fit(X_train,y_train).score(X_test)))

(1, 2), min_df=2: 70273
(1, 2), min_df=2 score: 0.9520422745830721
(1, 2), min_df=2 test score: 0.9098310604518625
(1, 2), min_df=4: 33764
(1, 2), min_df=4 score: 0.9497980487054856
(1, 2), min_df=4 test score: 0.909220435579076
```

Therefore, increase the value of min_df, the number of feature will decrease; the score will decrease.

5.0.2 impact of stop-words with n-grams

```
In [31]: pipe = make_pipeline(TfidfVectorizer(ngram_range=(1, 2), min_df=4), LogisticRegression)
         tfidf = TfidfVectorizer(ngram_range=(1, 2), min_df=4)
         tfidf.fit(X_train,y_train)
         print("(1, 2), min_df=4: {}".format(len(tfidf.vocabulary_)))
```

```

print('(1, 2), min_df=4 score: {}'.format(np.mean(cross_val_score(pipe, X_train, y_train, cv=5, scoring='f1'))))
print('(1, 2), min_df=4 test score: {}'.format(pipe.fit(X_train, y_train).score(X_test, y_test)))

pipe = make_pipeline(TfidfVectorizer(ngram_range=(1, 2), min_df=4, stop_words="english"))
tfidf = TfidfVectorizer(ngram_range=(1, 2), min_df=4, stop_words="english")
tfidf.fit(X_train, y_train)
print("(1, 2), stopwords, min_df=4 score: {}".format(len(tfidf.vocabulary_)))
print('(1, 2), stopwords, min_df=4 score: {}'.format(np.mean(cross_val_score(pipe, X_train, y_train, cv=5, scoring='f1'))))
print('(1, 2), stopwords, min_df=4 test score: {}'.format(pipe.fit(X_train, y_train).score(X_test, y_test)))

```

```

(1, 2), min_df=4: 33764
(1, 2), min_df=4 score: 0.9535003525255963
(1, 2), min_df=4 test score: 0.892937105638103
(1, 2), stopwords, min_df=4 score: 20623
(1, 2), stopwords, min_df=4 score: 0.94595407251938
(1, 2), stopwords, min_df=4 test score: 0.8927335640138409

```

Therefore, apply stop words, the number of feature will decrease; the score will also decrease.

6 Task 4

From task3, we know, when using (1,2) grams and min_df=4 the model has a best performance.

```

In [17]: from sklearn.svm import LinearSVC
         from sklearn.linear_model import RidgeClassifier

In [13]: X_train = train['Title'].map(str) + ' ' + train['Review']
         X_test = test['Title'].map(str) + ' ' + test['Review']

In [16]: pipe = make_pipeline(TfidfVectorizer(ngram_range=(1, 2), min_df=4), LogisticRegression())
         print('L1 cv score: {}'.format(np.mean(cross_val_score(pipe, X_train, y_train, cv=5, scoring='f1'))))
         print('L1 test score: {}'.format(pipe.fit(X_train, y_train).score(X_test, y_test)))

L1 cv score: 0.9448186145300251
L1 test score: 0.8966008548748219

In [18]: pipe = make_pipeline(TfidfVectorizer(ngram_range=(1, 2), min_df=4), RidgeClassifier())
         print('L2 cv score: {}'.format(np.mean(cross_val_score(pipe, X_train, y_train, cv=5, scoring='f1'))))
         print('L2 test score: {}'.format(pipe.fit(X_train, y_train).score(X_test, y_test)))

L2 cv score: 0.9537612594021473
L2 test score: 0.9102381437003867

In [19]: pipe = make_pipeline(TfidfVectorizer(ngram_range=(1, 2), min_df=4), LinearSVC())
         print('LinearSVC cv score: {}'.format(np.mean(cross_val_score(pipe, X_train, y_train, cv=5, scoring='f1'))))
         print('LinearSVC test score: {}'.format(pipe.fit(X_train, y_train).score(X_test, y_test)))

```

```
LinearSVC cv score: 0.9520345621531835
LinearSVC test score: 0.9118664766944841
```

From above models results:

- 1.L1 model has largest variance and bias.
- 2.LinearSVC has lowest variance and bias.
- 3.L2 model has a slightly worse performance than LinearSVC.

6.0.1 Other features

Beside using n-grams, we also could use:

- 1.sentiment score of the reviews and tiltes
- 2.Length of text
- 3.Number of out-of-vocabularly words
- 4.Presence / frequency of ALL CAPS
- 5.Lemmatization

In []: